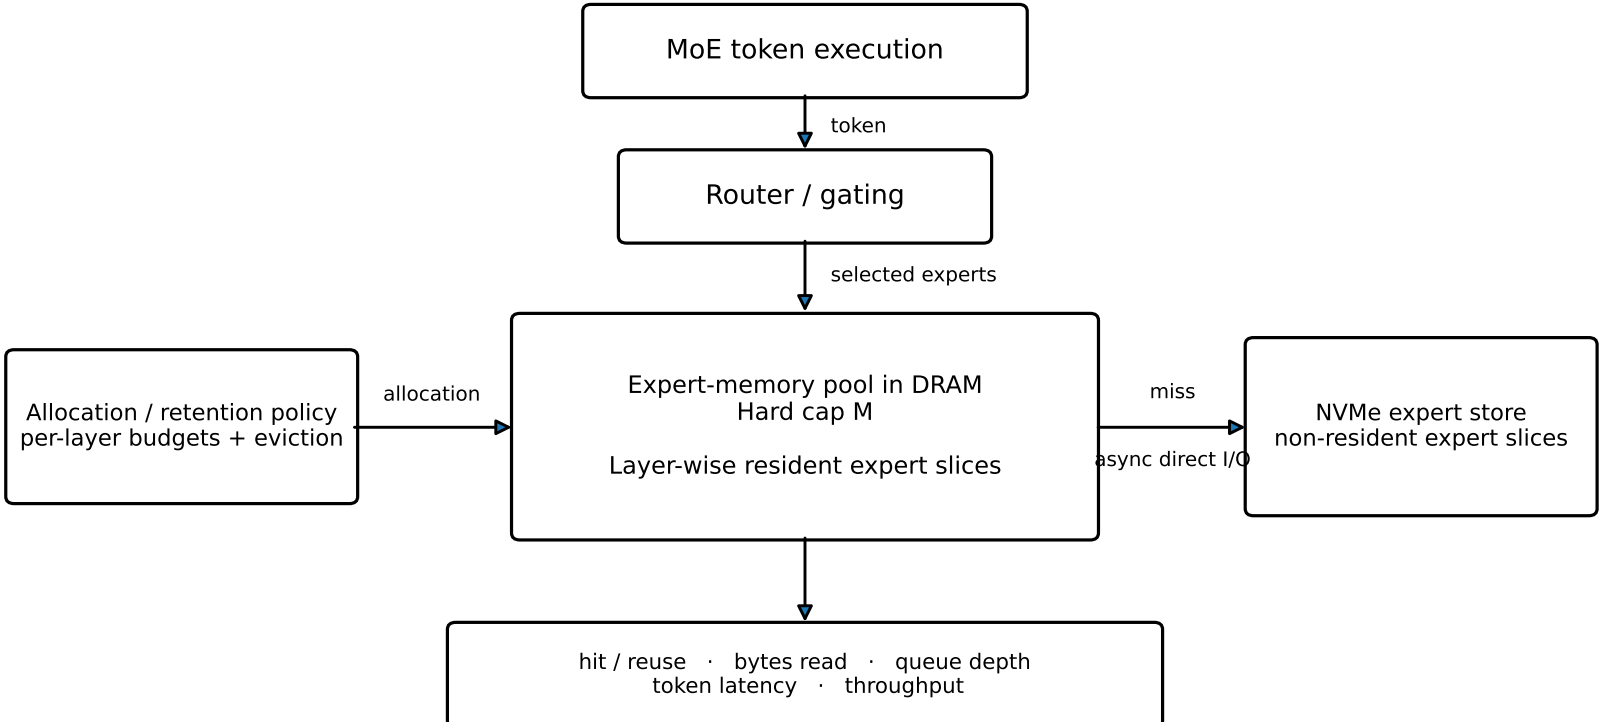
MoE token execution
token
Router / gating
selected experts
Allocation / retention policy
per-layer budgets + eviction
allocation
Expert-memory pool in DRAM
Hard cap M
Layer-wise resident expert slices
miss
async direct I/O
NVMe expert store
non-resident expert slices
hit / reuse · bytes read · queue depth
token latency · throughput